



DEVONEERS

RootRise

Redefining the Future of SME Growth ... AI &

RootRise Orchestrated

Breaking the Monolith, turning one stressed robot into a team of experts



RootRise



The Vision

We all agree on the vision:
RootRise isn't just a chatbot; it's a growth engine.

Currently, our content is Gold.

The methodology you wrote is world-class.

It's our 'Moat.'

But our delivery mechanism—the way the AI handles that content—is currently like one genius consultant trying to write a 100-page report in one hour.

Today, I'm showing you how we are going to architect this into a scalable 'Digital Factory.'



The Current Risk – "The Monolith"



"Risk: Hallucinations, Slow Speed, Inconsistent Results."

Right now, our backend uses what engineers call a '**Monolith.**'

We feed the AI one massive set of instructions—
Leadership, Ops, Finance, Export—all at once.

This is dangerous. It's like hiring a worker and giving them
a 500-page manual and saying 'Memorize this
and build the car.'

They will forget things. They will hallucinate.

And every time a client clicks 'Generate,' it costs us
money to send that huge manual.

We can't scale this to 4,000 SMEs.



The Solution - "Orchestration"



THE ROUTER

Our robot whose only job is to figure out exactly what the user needs.



This avoids confusion, overload, and hallucinations—because the model is no longer doing everything at once.

(The Assembly Line):

- Step 1: "The Router" (decides what the problem is).
- Step 2: "Specialized Agents"
- Step 3: "The Manager" (synthesizes the report).

Re-architecting the system to work like a factory line. Instead of one big bot, we will have **specialized agents**.

1. The **Ops Agent** only looks at production data.
2. The **Finance Agent** only looks at the numbers.
3. The **Manager Agent** combines their work. This means deeper insights, zero confusion, and faster results. The 'Logic' (our prompts) stays the same; we are just organizing who does what.

SPECIALIZED AGENTS

FINANCE BOT

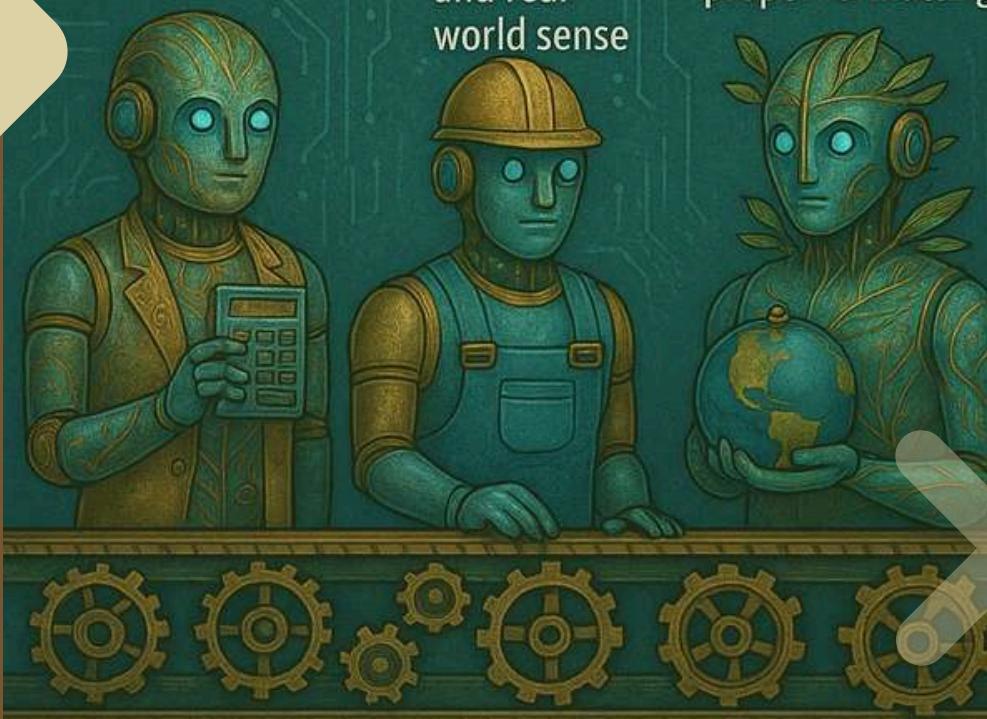
Crunches numbers
Builds budgets and forecasts
Validates financial logic

OPS BOT

Handles processes, operations, workflows
Ensures feasibility and real-world sense

EXPORT BOT

Manages trade rules, documentation, logistics
Ensures global compliance and proper formatting



Each agent receives only the part of the problem it's designed to handle—meaning fewer mistakes and faster results.



DEVONEERS

STEP 3 – THE MANAGER

- Collects outputs from all agents
- Cleans, merges and unifies everything
- Produces a consistent, human-ready final result
- Ensures accuracy, coherence, and structure

A coordinated team → faster, smarter, more accurate results. A factory assembly line – but for intelligence.



The Immediate Fix – The "Diagnostic Bank"

Regarding the Diagnostic Template—right now, we ask the AI to invent the questions every time a user logs in.

That's slow and risky.

We have to change this.

We will use our prompt to build a Master Question Vault in our database. When a user logs in, the system simply retrieves the perfect questions instantly. It's faster, cheaper, and 100% consistent for benchmarking.

THE IMMEDIATE FIX – THE “DIAGNOSTIC BANK”



VAULT



BEFORE
AI GENERATING
QUESTIONS
ON THE FLY
(SLOW/UNPREDICTABLE)



AFTER
AI SELECTING
QUESTIONS FROM
THE “VAULT”
(FAST/CONSISTENT)

Approach - "Building the Plane While Flying"

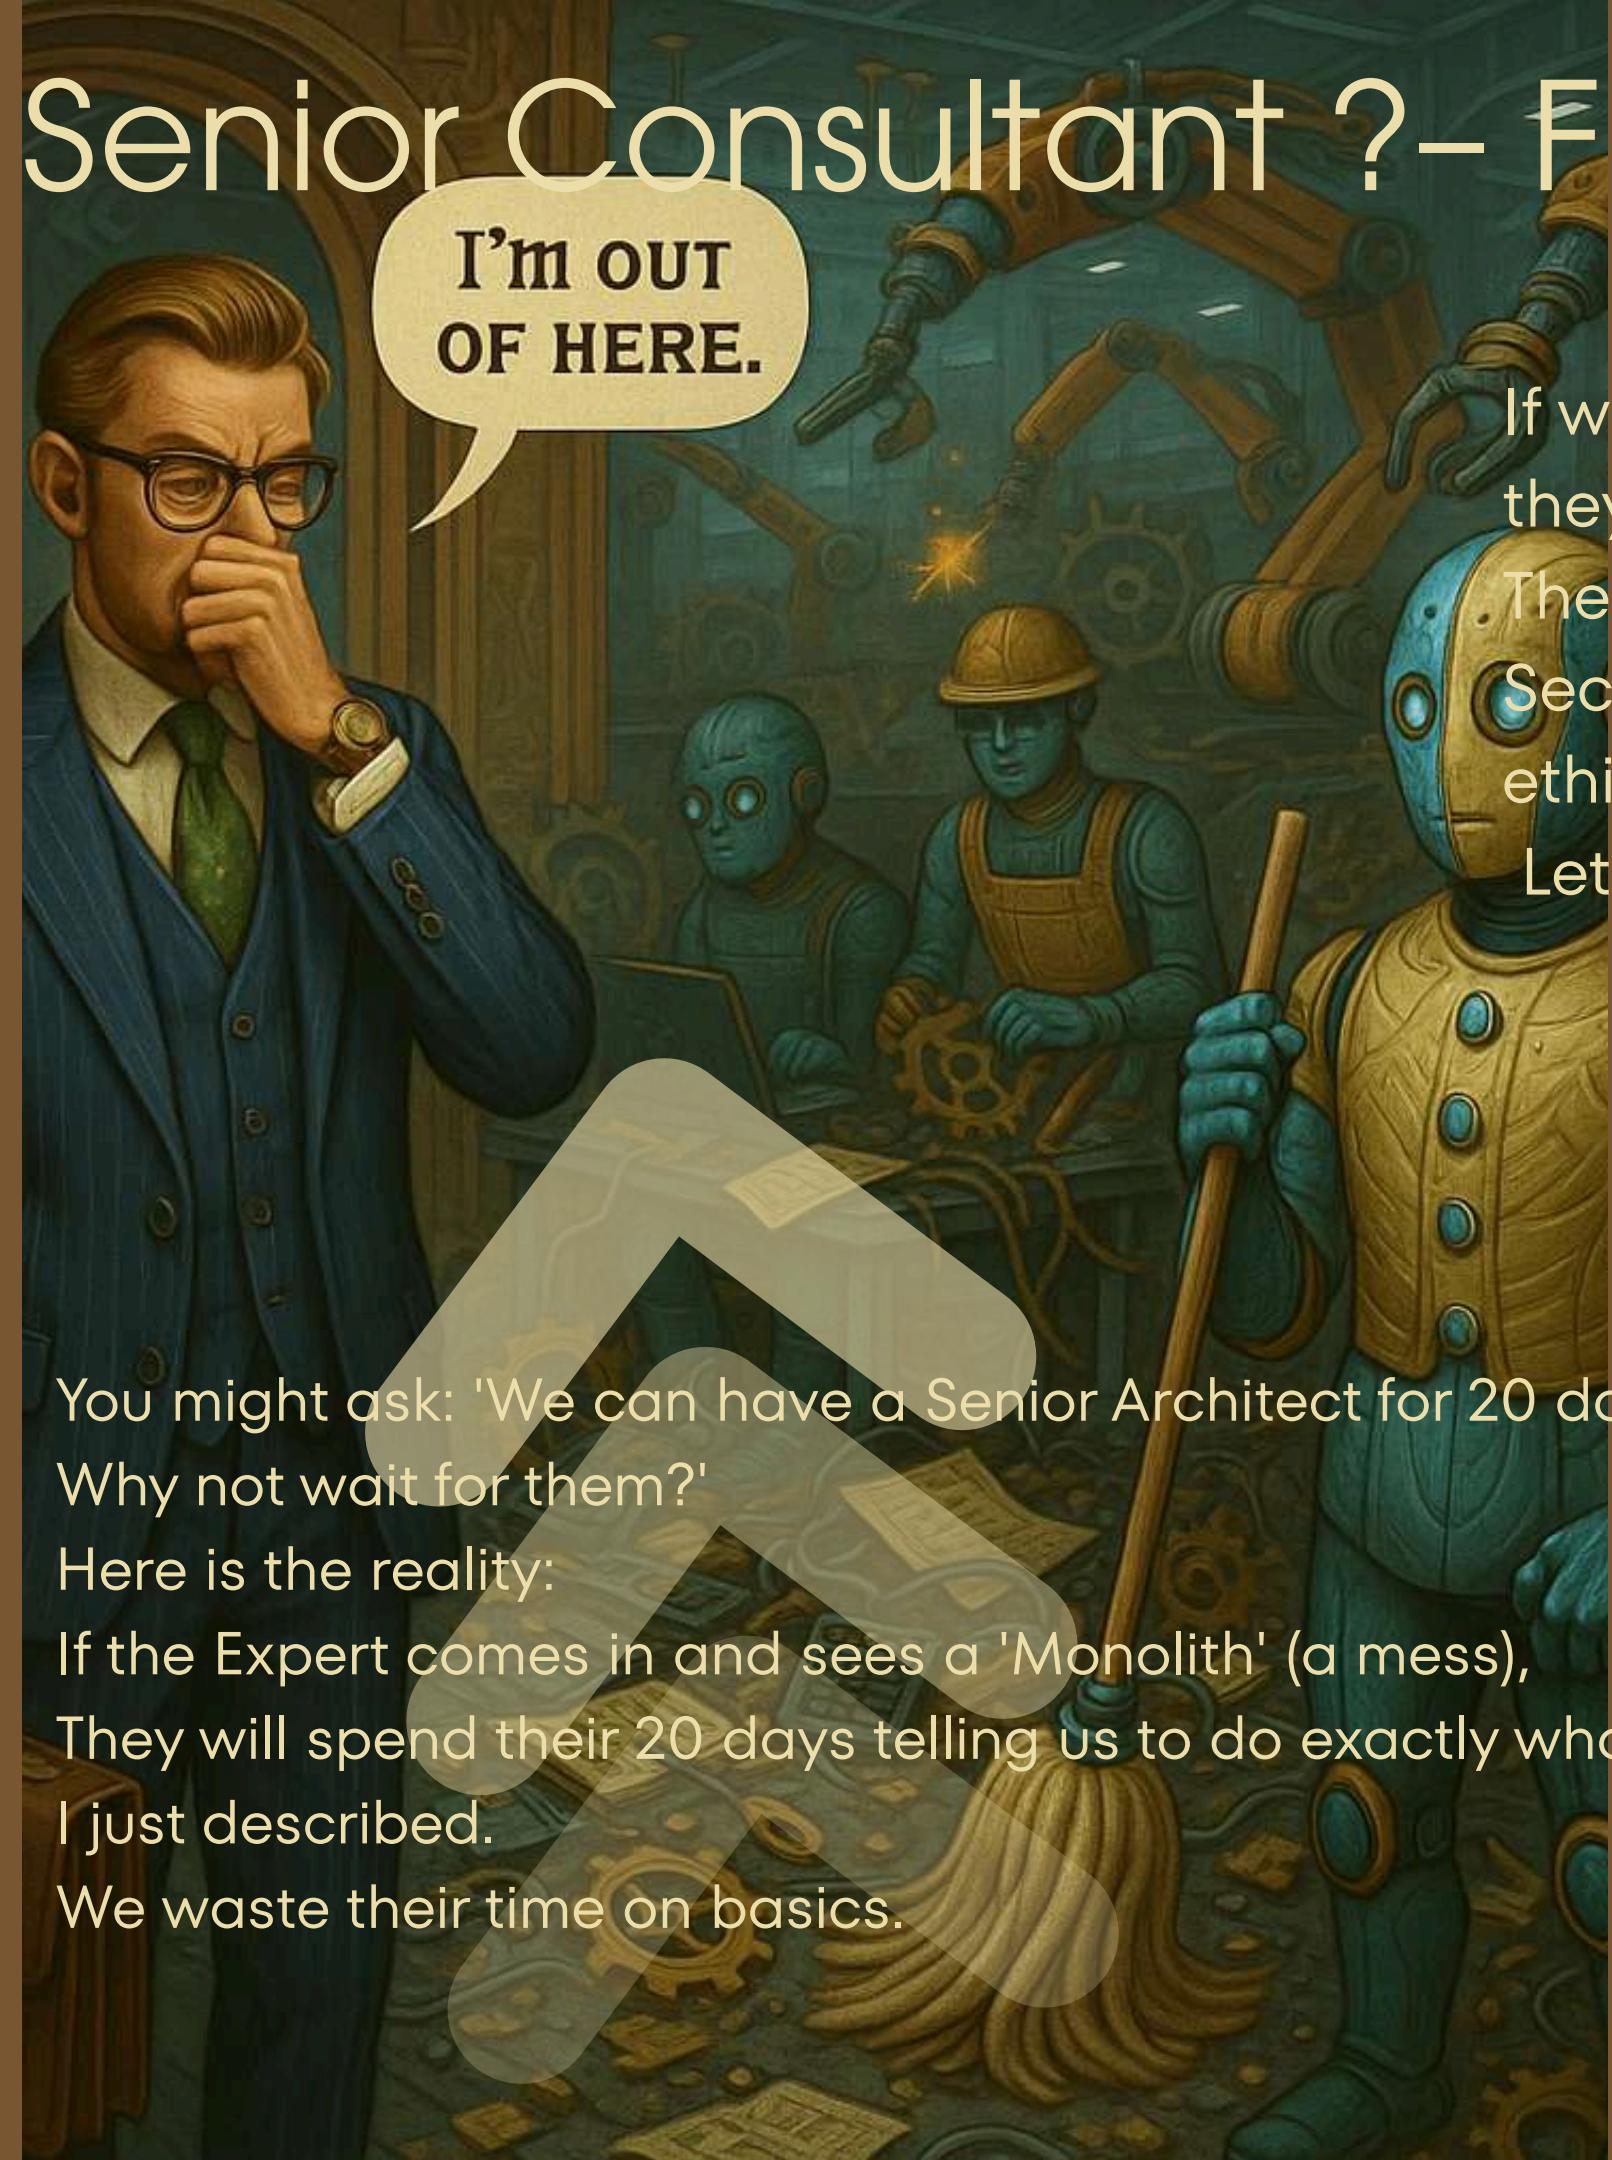
Diving deep into the advanced architecture required to build this.
Without waiting to finish 'studies' to fix the platform, applying the fixes live.
Every new architectural pattern learned (like Agent Handoffs),
will be implemented immediately.
This means the platform gets smarter every week.
We aren't waiting for a 'Big Bang' launch; we are iterating.



CONTINUOUS INTEGRATION



Senior Consultant ?- Fix It Now?



I'M OUT
OF HERE.

If we fix the architecture now, when the Expert arrives,
they will see a working factory.

Then, they can spend their 20 days on High-Level Strategy:
Security hardening, scaling to 4000 users, and advanced AI
ethics.

Let's use the Expert for Optimization, not Cleanup."

You might ask: 'We can have a Senior Architect for 20 days.

Why not wait for them?'

Here is the reality:

If the Expert comes in and sees a 'Monolith' (a mess),
They will spend their 20 days telling us to do exactly what
I just described.

We waste their time on basics.



The "Human-in-the-Loop" Guarantee

- Green: Safe recommendations.
- Red: Complex decisions -> Diverts to Human (**&I**)

Finally, this new architecture allows us to insert the **validation gates** we promised investors.

In a Monolith, it's hard to stop the AI.

In an assembly line, we can put a 'Quality Control Station' between the Finance Agent and the Final Report. The AI does the heavy lifting, but we (the experts) verify the output before the client ever sees it.

That is the true '**&I**' factor."

THE 'HUMAN-IN-THE-LOOP' GUARANTEE



GREEN:
Safe recommendations.



RED:
Complex decisions



